

Through-focal inversion of LEEM intensities with the Fourier-optics kernel

Wilson Lau

26 August 2026

Abstract

We present a through-focal inversion of low-energy electron microscopy (LEEM) intensities for the object exit wave $\psi_0 = \sigma e^{i\phi}$, based on the Fourier-optics (FO) kernel of Yu, Lau and Altman [11]. Prior FO studies compared predicted and measured contrast for prescribed objects; they did not recover ψ_0 . We invert that map in two stages. Stage 1 is regularized multi-defocus inversion on the contrast-transfer slice $R_{\text{FO}}(\mathbf{q}, 0, \Delta z)$, which fills zeros of a single-defocus contrast-transfer function (CTF) and is the unique vacuum 2×2 Tikhonov solve for $\alpha > 0$. For $\phi \ll 1$ the large-phase estimator is a mixed transmission-cross-coefficient (TCC)–Born homotopy: rank- M TCC apply (exact rank-1 autocorrelation under perfect coherence), remainder-corrected Born on the same Fourier-diagonal 2×2 , per-bin skip $t \in \{0, 1\}$, and optional algebraic Armijo on the exact quartic line energy. A one-dimensional sinusoid reduces the second stage to Jacobi–Anger fitting of the aperture-truncated Bessel ladder, not a two-dimensional dispatcher. Uniqueness of the regularized Fourier-diagonal estimator, the quadratic homotopy, the TCC apply and the quartic line energy are machine-checked in Lean 4 against that kernel. The catalogue of every competing inverse examined, with Pick / Not-pick reasons, is the companion selection note [28]. This work extends quantitative evaluation of LEEM contrast from forward matching of a focal series to recovery of the object that produced it.

1 Introduction

Image interpretation in LEEM [1, 2, 3, 4] requires a quantitative account of how the cathode lens, contrast aperture, source size and energy spread map the object exit wave to the recorded intensity. Aberration-corrected instruments [5, 6] make that map sharper, but they do not remove the distinction between two standard calculations. The contrast-transfer function (CTF) evaluates the inverse transforms of ΨH and $\Psi^* H^*$ separately and is computationally efficient [8, 9]. Fourier optics (FO) retains inseparable mixing of spatial-frequency pairs $(\mathbf{q}, \mathbf{q}')$ under partial coherence [7, 10, 11]. The mathematically rigorous FO calculation gains an advantage under conditions, particularly defocus, in which partial coherence compromises the validity of the CTF intensity [11]. Yu, Locatelli and Altman [10] showed that this distinction is not academic: CTF modelling fails for strongly rippled graphene, whereas FO reproduces the measured through-focal contrast.

Those studies prescribe an object, evaluate the forward map, and compare predicted and measured contrast versus defocus. They do not invert. The only demonstrated 2D experimental LEEM exit-wave reconstruction is the linear focal-series restoration of Duden *et al.* [12], in the tradition of Schiske [14, 15], assuming a strong specular beam and a CTF-style model. Here the same FO forward map is inverted: given a through-focal intensity series $I(\mathbf{r}, \Delta z_k)$, $k = 1, \dots, K$, we recover ψ_0 within the aperture cutoff, up to a global phase. The Lean identities of the companion note [27] are the prerequisite: the discrete inverse is checked against that kernel, not against a substitute propagator.

A single defocus underdetermines a general complex object. The CTF factor may vanish on interior rings; the contrast aperture discards $|\mathbf{q}| > q_{\text{ap}}$; intensity is invariant under a global phase; and one complex measurement cannot identify the conjugate pair $(X(\mathbf{q}), X(-\mathbf{q}))$. A focal series with non-coincident zeros restores those modes that the aperture admits. We compare several inverses and recommend a two-stage estimator that uses the Yu 2019 kernel on the CTF slice $R_{\text{FO}}(\mathbf{q}, 0, \Delta z)$ as initialization and as the quiet-bin branch, and a mixed TCC–Born homotopy on the full bilinear residual for loud bins. Selected algebraic properties of the discrete estimator are machine-checked in Lean 4 [23] against that kernel (Section 4 and Appendix A). Every method examined is recorded in [28].

2 Image formation

The object is the complex exit wave

$$\psi_0(\mathbf{r}) = \sigma(\mathbf{r}) e^{i\phi(\mathbf{r})}, \quad (1)$$

with ϕ in radians. For a topographic height A in reflection the kinematic phase is

$$\phi = \frac{4\pi A}{\lambda_0}, \quad (2)$$

where λ_0 is the wavelength at the sample. At $E_0 = 10$ eV, $\lambda_0 \simeq 0.39$ nm, so $A = 0.1$ nm already gives $\phi \sim 3$, outside the weak-phase regime occupied by graphene ripples and atomic steps [10]. Lens optics use the wavelength λ after acceleration to column energy E (typically 15 kV): spatial frequency is $\mathbf{q} = \boldsymbol{\alpha}/\lambda$ [9, 8]. The reconstruction kernel below is written in λ ; conversion of reconstructed phase to height uses λ_0 .

With $\Psi = \mathcal{F}[\psi_0]$, the FO intensity is bilinear in the object spectrum [13, 7, 11],

$$I(\mathbf{r}, \Delta z) = \iint \Psi(\mathbf{q}) \Psi^*(\mathbf{q}') R_{\text{FO}}(\mathbf{q}, \mathbf{q}', \Delta z) e^{2\pi i(\mathbf{q}-\mathbf{q}')\cdot\mathbf{r}} d^2q d^2q'. \quad (3)$$

The kernel is supported on the contrast aperture $|\mathbf{q}|, |\mathbf{q}'| \leq q_{\text{ap}}$ and factors as $R_{\text{FO}} = R_0 E_S E_{C,\text{tot}}$, with closed-form spatial and chromatic envelopes derived in [11] and proved in [27]. The CTF kernel agrees with R_{FO} on the slice $\mathbf{q}' = \mathbf{0}$,

$$R_{\text{FO}}(\mathbf{q}, 0, \Delta z) = R_{\text{CTF}}(\mathbf{q}, \Delta z). \quad (4)$$

The axis $\mathbf{q} = \mathbf{0}$ is the conjugate branch $R_{\text{FO}}(0, -\mathbf{q})$, not the CTF. Wave aberration χ is measured in waves, with transfer factor $W = e^{2\pi i\chi}$ [9]. Defocus enters χ_S as $+\frac{1}{2}\Delta z\lambda q^2$ in the convention of this work; mixing the opposite sign used in some simulators interchanges overfocus and underfocus.

Intensity is invariant under $\psi_0 \mapsto e^{i\theta}\psi_0$. Modes with $|\mathbf{q}| > q_{\text{ap}}$ are identically invisible in every defocus. The linearization of (3) about a strong specular beam therefore recovers ψ_0 only up to a global phase (fixed here by taking the zero-frequency component real and positive) and only inside the aperture.

3 Through-focal inversion

Let N denote the number of image pixels and K the number of defoci; a typical stack is $N = 512^2$ and $K \sim 11$ –21. Regularization of an ill-posed filter follows Tikhonov [17] and Wiener [18]: a positive floor α in the denominator prevents division by a vanished CTF at the price of a bias of order α .

3.1 Candidate methods

Single-image Wiener deconvolution applies one Fourier transform and divides by $R_{\text{CTF}}(\mathbf{q}, \Delta z)$. Interior CTF zeros leave undetermined modes, and one plane does not identify a general complex object.

Transport of intensity [21, 22] recovers phase from two or three closely spaced defoci under a paraxial continuity equation. After acceleration, LEEM has $\lambda \sim 10$ pm, a hard aperture, spherical aberration and partial coherence. That PDE is not an inverse of (3) and has not been the method used on experimental LEEM graphene or steps.

Multi-defocus Schiske restoration [14, 12] transforms each plane and treats spatial frequencies independently. The slice multipliers are

$$h_k(\mathbf{q}) = R_{\text{FO}}(\mathbf{q}, 0, \Delta z_k). \quad (5)$$

For a single unknown amplitude $X(\mathbf{q})$ the regularized estimator is

$$\hat{x} = \frac{\sum_k \overline{h_k} y_k}{\alpha + \sum_k |h_k|^2}. \quad (6)$$

Linearization about vacuum couples $X(\mathbf{q})$ with $X(-\mathbf{q})$; the corresponding 2×2 regularized system uses the same α and is not (6). If one defocus has $h_k = 0$, the remaining planes can still determine the bin. The cost is $K+1$ discrete Fourier transforms plus $O(KN)$ arithmetic. This is Duden’s algorithm with the Yu 2019 envelopes, and it coincides with one Gauss–Newton step from vacuum on (3).

Gerchberg–Saxton iteration [19] and hybrid input–output [20] replace the computed modulus by the measured intensity and propagate with a coherent lens model. Under partial coherence the intensity constraint is the quadratic set of (3). Naive modulus replacement is therefore the coherent CTF model and is biased when FO and CTF disagree off axis [10]. Projection onto the FO intensity is a slower, undamped rewrite of nonlinear least squares.

Nonlinear least squares minimizes the bilinear FO residual of the stack. Each Gauss–Newton step uses the pair kernel, so it sees the mixing dropped by (6). That estimator is FO-faithful for $\phi \ll 1$, but a dense pair apply is KN^2 and is not the recommended mix. The large-phase 2D estimator is the mixed TCC–Born homotopy of Section 3.2: the same residual evaluated by a rank- M TCC (or a coherent rank-1 autocorrelation) and a Fourier-diagonal 2×2 Born step, with per-bin skip $t \in \{0, 1\}$. If the aberrations themselves are unknown, the same idea stacked over extra defoci is the maximum-likelihood (MAL) reconstruction of Coene, Thust, Op de Beeck and Van Dyck [16].

Method	Typical cost	Fills CTF zeros	Faithful for large ϕ
Single-image Wiener	one transform	no	no
Transport of intensity	one transform	n/a (wrong PDE)	no
Multi-defocus regularized	$K+1$ trans-forms	yes ($K \geq 2$)	linearized only
Gerchberg–Saxton / HIO	many transforms	if $K \geq 2$	only if the projector is FO
Dense nonlinear least squares on FO	pair sums KN^2	yes	yes (not cheap)
Mixed TCC–Born homotopy	$O(KMN \log N)$	yes	yes (2D pick)
Forgetful Kaczmarz ($K \geq 2$)	$O(K)$ per bin, wrong min	yes if $K \geq 2$	linearized only, not Tikhonov
Jacobi–Anger fit (1D sinusoid)	$O(M^2)$	n/a (parametric)	yes (1D only)

3.2 Recommended estimator

For a general 2D experimental stack the fastest method that still fills CTF zeros in the *linearized* regime is multi-defocus regularized inversion on the slice (4). It reads all KN Fourier samples,

fills interior zeros of $\sin(2\pi\chi_S)$ by stacking defoci, and has a unique minimizer for every $\alpha > 0$ (Theorem A.1). That map is also stage 1 of the large-phase mix and the quiet-bin branch.

For $\phi \not\ll 1$ the FO-faithful 2D estimator initializes that reconstruction and applies a mixed Born homotopy: the bilinear remainder is evaluated with a rank- M TCC (Hopkins) expansion, or exactly as an autocorrelation when the source is a point, optionally with a diagonal leftover (LR+D), and the same Fourier-diagonal 2×2 system is solved on the corrected residual. Skip is per-bin: quiet bins keep the stage-1 residual ($t = 0$), loud bins take full Born ($t = 1$). Rank policy at $N = 128$: $M = 1$ under perfect coherence or when the exact quartic line search is used, otherwise $M \leq 8$. An algebraic Armijo step on that quartic is the FO-faithful 1D mix after a Born direction (a unit Gauss–Newton step is not a critical point of the quartic, and the Gauss–Newton Hessian drops the remainder coupling). Forgetful Kaczmarz is not the unique Tikhonov minimizer for $K \geq 2$. Dense Gauss–Newton about a generic background remains FO-faithful but is not Fourier-diagonal and is not $O(KN \log N)$. MAL applies if aberrations are unknown. A global omit-stage-2 certificate exists when every bin is quiet; the numerical reading $\max|\phi| \lesssim 0.3$ at typical LEEM signal-to-noise ($A \lesssim 0.01$ nm at $E_0 = 10$ eV) is not part of the machine-checked fragment. The linearization never sees off-axis pairs $(\mathbf{q}, \mathbf{q}')$ at which FO and CTF disagree; that omitted term is the quadratic remainder of (3) about vacuum (Theorem A.4) and is the whole image for a strong-phase sinusoid [10]. Agreement of the slice (4) with a CTF calculation is therefore not a sufficient skip criterion.

For the one-dimensional sinusoid of period Λ used in [10, 11] to exhibit CTF failure, the unknown reduces to an amplitude φ and a shift. The aperture retains $M = 2\lfloor q_{\text{ap}}\Lambda \rfloor + 1$ modes. The FO image is then a double sum of M^2 Jacobi–Anger terms,

$$e^{i\varphi \sin(2\pi x/\Lambda)} = \sum_{n \in \mathbb{Z}} J_n(\varphi) e^{2\pi i n x/\Lambda}, \quad (7)$$

and least squares on φ with $J'_n = \frac{1}{2}(J_{n-1} - J_{n+1})$ is $O(M^2)$. A general 2D micrograph has no Bessel ladder; the nonlinear stage is the mixed TCC–Born homotopy in the aperture basis, not a dense pair-sum Gauss–Newton loop.

3.3 Discrete procedure

1. Rigid alignment (informal) and mean-subtracted discrete Fourier transforms of the stack; the DC bin carries total flux ($R_{\text{FO}}(0, 0) = 1$).
2. For every $|\mathbf{q}| \leq q_{\text{ap}}$ and recorded defocus, evaluate the slice (5) from the Yu 2019 envelopes. Default $K = 2$ independent defoci; coefficients with $|\mathbf{q}| > q_{\text{ap}}$ are set to zero.
3. Stage 1: the conjugate-pair 2×2 Tikhonov solve at each bin with fixed $\alpha > 0$ (`stage1Pair2`, `tikhonovXhat2`). The Hermitian partner at $-\mathbf{q}$ is recovered from an odd `qmap` when the data are Hermitian (`mixed2D_conj_partner`).
4. Per-bin skip: if every defocus remainder at that bin is at most η , keep $t = 0$ (stage-1 residual); otherwise $t = 1$ (`remainderWeight`, `mixed2DMix`).
5. Loud bins: apply $\hat{I}(\delta)$ with rank- M TCC (`ihat_tcc`, `recommendTccRank`; exact rank-1 under `PerfectCoherence`), form the remainder-corrected residual, and re-solve the vacuum 2×2 (`mixedBinStep`). Rank policy at $N = 128$: $M = 1$ if perfectly coherent or if a quartic line search is used, else $M \leq 8$.
6. Optional algebraic Armijo along that Born direction (`lineDescentStep` on `lineFid`).
7. Inverse transform yields $\sigma = |\hat{\psi}|$ and $\phi = \arg(\hat{\psi})$ in the vacuum gauge.

Experimental lens current is mapped onto the defocus convention of Section 2 by one fitted affine transform (informal).

4 Machine-checked properties of the discrete estimator

Selected algebraic claims for the discrete linear estimator and the mixed TCC–Born / quartic-line stage are machine-checked in Lean 4 [23, 24] against the FO kernel encoded from Yu *et al.* [11, 27]. The statements are finite-dimensional (finite defoci and a finite frequency grid). The discrepancy n_k that enters the error identity is deterministic (linearization remainder, instrument noise and discretization); no probability space is introduced. Full statements are collected in [Appendix A](#).

Stage 1 uses (4), not a substitute propagator; for $\phi \ll 1$ the estimator is the mixed TCC–Born homotopy of [Section 3.3](#), not a dense pair-sum Gauss–Newton loop. For $\alpha > 0$ the regularized energy at each frequency bin has a unique global minimizer (6) ([Theorem A.1](#)). If $y_k = h_k x^* + n_k$, the estimation error is exactly $(\sum_k \overline{h_k} n_k - \alpha x^*)/D$ with $D = \alpha + \sum_k |h_k|^2$ ([Theorem A.2](#)); a sharp triangle bound follows ([Theorem A.3](#)). Frequencies outside the aperture reconstruct as zero at every defocus; a global phase is invisible; one measurement does not identify $(X(\mathbf{q}), X(-\mathbf{q}))$; and $\sin(2\pi\chi_S)$ vanishes on interior rings of a pure-defocus column ($C_3 = C_5 = 0$, $\Delta z \neq 0$, large enough aperture), which is why $K \geq 2$ is required ([Theorem A.6](#)). The multi-defocus slice map is the derivative of bilinear FO at vacuum; the neglected term is homogeneous of degree two ([Theorem A.4](#)). When that remainder lies above the noise, the nonlinear stage is the mixed TCC–Born homotopy with per-bin skip (`remainderWeight`, `mixed2DMix`); the numerical criterion $\max|\phi| \lesssim 0.3$ and the global `stage2Skip` certificate are not a second skip policy.

A modelled operation count, treating a discrete Fourier transform as $N \log_2 N$ operations, shows that the linear stage is cheaper than a dense pair apply already at $N = 128$, and that a rank- $M \leq 8$ TCC Born step remains cheaper than one dense apply at that grid. Existence of a fast Fourier transform, wall-clock times, Gaussian or Poisson statistics, Cramér–Rao bounds, Born/Gauss–Newton convergence, Cardano roots of the line cubic, image alignment and the current-to-defocus map are not claimed.

5 Conclusions

We invert the FO image-formation map of Yu, Lau and Altman [11] for a 2D through-focal intensity series. Regularized multi-defocus inversion on the CTF slice fills the zeros of a single defocus, is unique for every positive regularization, and is first-order consistent with bilinear FO: it is stage 1 of the mix and the quiet-bin branch. For $\phi \ll 1$ the 2D estimator is the mixed TCC–Born homotopy: rank- M TCC (or coherent rank-1) apply, remainder-corrected Born on the Fourier-diagonal 2×2 , per-bin skip $t \in \{0, 1\}$, and optional quartic Armijo [28]. For the one-dimensional sinusoid that exhibits CTF failure the second stage collapses to a Bessel fit. These results broaden quantitative LEEM contrast analysis from forward matching of a focal series to recovery of the object, with the discrete estimator verified against the same kernel used to predict the images.

A Formal claims

The statements below are finite-dimensional. No existence theorem for the fast Fourier transform is claimed.

Theorem A.1 (Unique minimizer). *For $\alpha > 0$ the energy $J(x) = \sum_k \|h_k x - y_k\|^2 + \alpha \|x\|^2$ has a unique global minimizer (6). Completing the square: $J(x) = J(\hat{x}) + D \|x - \hat{x}\|^2$ with $D = \alpha + \sum |h_k|^2 > 0$.*

Lean: `tikhonovJ_min`, `tikhonov_unique`, `tikhonovJ_eq_shift`. The same identity gives uniqueness at $\alpha = 0$ whenever $\sum |h_k|^2 > 0$.

Theorem A.2 (Bias–noise identity). *If $y_k = h_k x^* + n_k$ and $D \neq 0$, then*

$$\hat{x} - x^* = \frac{\sum_k \overline{h_k} n_k - \alpha x^*}{D}. \quad (8)$$

Lean: `tikhonov_error`. Taking $\mathbb{E}[n] = 0$ is informal (no probability space).

Theorem A.3 (Triangle bound, sharp). *If $\alpha \geq 0$ and $D > 0$,*

$$|\hat{x} - x^*| \leq \frac{\sum_k |h_k| |n_k| + \alpha |x^*|}{D}. \quad (9)$$

Equality holds for one measurement, $h = n = 1$, $x^ = -1$, $\alpha = 1$.*

Lean: `tikhonov_error_bound`, `tikhonov_error_bound_sharp`.

Theorem A.4 (Quadratic remainder of bilinear FO). *With vacuum spectrum $1_{\mathbf{q}=0}$ and increment δ ,*

$$\hat{I}(\text{vac} + \delta) = \hat{I}(\text{vac}) + DI_{\text{vac}}[\delta] + \hat{I}(\delta), \quad (10)$$

where $DI_{\text{vac}}[\delta](\xi) = R_{\text{FO}}(\xi, 0) \delta(\xi) + R_{\text{FO}}(0, -\xi) \overline{\delta(-\xi)}$ and $|\hat{I}(\delta)(\xi)| \leq (\sum |R_{\text{FO}}|) (\sum |\delta|)^2$.

Lean: `ihat_add`, `ihatJac_vacuum`, `ihat_quadratic_remainder`, `ihat_bound`. The dropped term is exactly bilinear: linearized inversion is biased by $O(\|\delta\|^2)$, which is the whole image for a strong-phase sinusoid [10]. There is no cubic remainder along the real ray (`ihat_homotopy_cubic_zero`).

Theorem A.5 (Per-bin skip). *The recommended remainder weight is 0 if every defocus remainder at the bin is at most η , and 1 otherwise (`remainderWeight_eq_zero_or_one`). The $n = 0$ iterate of the recommended mix is the 2D stage-1 map (`mixed2DMix_zero_eq_stage1Pair2`).*

Modelled TCC cost. At $N = 128$, a rank- $M \leq 8$ TCC apply is strictly cheaper than dense KN^2 (`tccApplyCost_lt_dense_128`); $M = 9$ remains cheaper; $M \geq 10$ is not. The Born cap stays $M \leq 8$. Rank-1 quartic line search is cheaper than dense at that grid; rank-8 line search is not.

Theorem A.6 (Interior CTF zeros of a pure-defocus column). *$\sin(2\pi\chi_S) = 0$ whenever $2\chi_S \in \mathbb{Z}$. For $C_3 = C_5 = 0$, $\Delta z \neq 0$, and a large enough aperture, there exists nonzero q inside the disk with that vanishing (`exists_interior_weakPhase_zero`). A single defocus is therefore blind at those rings; $K \geq 2$ with incommensurate Δz typically fills them. This is a pure-defocus hypothesis, not a column that still has spherical aberration.*

The 2×2 Gram determinant of the vacuum Jacobian is positive for $\alpha > 0$ (`gramDet_pos`), so the conjugate-pair normal equations are invertible even though one measurement is not ([Section 1](#)). The Cramer closed form is the unique minimizer (`tikhonovXhat2`, `tikhonov2_unique`). One vacuum Gauss–Newton residual is that 2×2 model, hence the closed-form 2D step is `stage1Pair2` (`vacuumGN_eq_stage1Pair2` in `Plane2.lean`). The 1D analogue is `stage1Pair` / `vacuumGN_eq_stage1Pair` in `Pipeline.lean`.

Modelled cost. Set $\text{dftCost}(N) := N \log_2 N$ and

$$\text{reconstructCost}(K, N) = (K + 1) N \log_2 N + N(8K + 4). \quad (11)$$

For $K \geq 1$ and $N \geq 2$, $\text{reconstructCost}(K, N) \leq 14 K N \log_2 N$. At $N = 128$ this is strictly cheaper than a dense bilinear apply KN^2 , for every $K \geq 1$ (`reconstructCost_le`, `reconstructCost_lt_dense_128`). The constant 14 is attained in the model at $K = 1$, $N = 2$.

Claim	Lean
Tikhonov unique min, closed form	LeemF0/Inverse/Tikhonov.lean
Bias–noise identity and sharp triangle bound	same
2×2 Cramer form, unique min	tikhonovXhat2, tikhonov2_unique
Cost model versus dense KN^2	reconstructCost_*
Aperture support of $R_{FO}(\cdot, 0)$	LeemF0/Inverse/LinearInverse.lean
Gauge, vacuum Jacobian, quadratic remainder	same
$K = 1$ not injective; pure-defocus CTF zeros	same
Hermitian kernel; DC $R_{FO}(0, 0) = 1$	R_F0_hermitian, R_F0_dc
Stage-1 map, vacuum GN glue, global skip certificate	LeemF0/Inverse/Pipeline.lean
2D kernel R_{FO2} , disk aperture	LeemF0/Forward/Kernel12.lean
2D stage-1 map <code>stage1Pair2</code>	LeemF0/Inverse/Plane2.lean
Rank-1 / TCC apply, hybrid cost	LeemF0/Inverse/LowRank.lean
Quadratic homotopy, Born/mixed2DMix skip	Homotopy.lean, Mix.lean; remainderWeight_eq_zero_or_on
Exact quartic line energy, <code>nlsJ_line</code>	LeemF0/Inverse/LineSearch.lean
TIE/GS/HIO/Rytov/Wiener/radial/JA rejects	exists_*_ne_* in Mix.lean
Mode count $M = 2\lfloor q_{ap}\Lambda \rfloor + 1$	LeemF0/Inverse/Modes.lean
FO ladder sample, <code>sinusoidJ</code> cost	rF0, sinusoidJ
Jacobi–Anger and Bessel decay	LeemF0/Forward/PhaseObject.lean
FFT exists / equals the continuum FT	informal
Gaussian/Poisson noise, Cramér–Rao	informal
Gauss–Newton local quadratic convergence	informal
Cardano roots of the line cubic; Banach FPT existence	informal
Nesterov / L-BFGS / Anderson / TV / RL rates	informal; see [28]
Image alignment, current-to-defocus map	informal
“Fastest among all conceivable algorithms”	only DFT versus dense bilinear
Full Pick / Not-pick catalogue	[28]

References

- [1] E. Bauer, *Low energy electron microscopy*, Rep. Prog. Phys. **57** (1994) 895–938.
- [2] E. Bauer, *Surface Microscopy with Low Energy Electrons* (Springer, New York, 2014).
- [3] R. M. Tromp, *Low-energy electron microscopy*, IBM J. Res. Dev. **44** (2000) 503–516.
- [4] M. S. Altman, *Trends in low energy electron microscopy*, J. Phys.: Condens. Matter **22** (2010) 084017.
- [5] R. M. Tromp, J. B. Hannon, A. W. Ellis, W. Wan, A. Berghaus, O. Schaff, *A new aberration-corrected, energy-filtered LEEM/PEEM instrument. I. Principles and design*, Ultramicroscopy **110** (2010) 852–861.
- [6] R. M. Tromp, J. B. Hannon, W. Wan, A. Berghaus, O. Schaff, *A new aberration-corrected, energy-filtered LEEM/PEEM instrument II. Operation and results*, Ultramicroscopy **127** (2013) 25–39.
- [7] A. B. Pang, Th. Müller, M. S. Altman, E. Bauer, *Fourier optics of image formation in LEEM*, J. Phys.: Condens. Matter **21** (2009) 314006.
- [8] S. M. Schramm, A. B. Pang, M. S. Altman, R. M. Tromp, *A contrast transfer function approach for image calculations in standard and aberration-corrected LEEM and PEEM*, Ultramicroscopy **115** (2012) 88–108.
- [9] R. M. Tromp, S. M. Schramm, *Optimization and stability of the contrast transfer function in aberration-corrected electron microscopy*, Ultramicroscopy **125** (2013) 72–80.

- [10] K. M. Yu, A. Locatelli, M. S. Altman, *Comparing Fourier optics and contrast transfer function modeling of image formation in low energy electron microscopy*, Ultramicroscopy **183** (2017) 109–116.
- [11] K. M. Yu, K. L. W. Lau, M. S. Altman, *Fourier optics of image formation in aberration-corrected LEEM*, Ultramicroscopy **200** (2019) 160–168.
- [12] T. Duden, A. Thust, C. Kumpf, F. S. Tautz, *Focal-series reconstruction in low-energy electron microscopy*, Microsc. Microanal. **20** (2014) 968–973.
- [13] J. W. Goodman, *Introduction to Fourier Optics*, 3rd ed. (Roberts and Company, Englewood, CO, 2005).
- [14] P. Schiske, *Zur Frage der Bildrekonstruktion durch Fokusreihen*, in *Electron Microscopy 1968* (Proc. 4th European Regional Conference, Rome), ed. D. S. Bocciarelli, vol. 1 (Tipografia Poliglotta Vaticana, Rome, 1968), pp. 145–146.
- [15] P. Schiske, *Image processing using additional statistical information about the object*, in *Image Processing and Computer-Aided Design in Electron Optics*, ed. P. W. Hawkes (Academic Press, London, 1973), pp. 82–90.
- [16] W. M. J. Coene, A. Thust, M. Op de Beeck, D. Van Dyck, *Maximum-likelihood method for focus-variation image reconstruction in high resolution transmission electron microscopy*, Ultramicroscopy **64** (1996) 109–135.
- [17] A. N. Tikhonov, *Solution of incorrectly formulated problems and the regularization method*, Sov. Math. Dokl. **4** (1963) 1035–1038.
- [18] N. Wiener, *Extrapolation, Interpolation, and Smoothing of Stationary Time Series* (MIT Press and Wiley, New York, 1949).
- [19] R. W. Gerchberg, W. O. Saxton, *A practical algorithm for the determination of the phase from image and diffraction plane pictures*, Optik **35** (1972) 237–246.
- [20] J. R. Fienup, *Phase retrieval algorithms: a comparison*, Appl. Opt. **21** (1982) 2758–2769.
- [21] M. R. Teague, *Deterministic phase retrieval: a Green’s function solution*, J. Opt. Soc. Am. **73** (1983) 1434–1441.
- [22] D. Paganin, K. A. Nugent, *Noninterferometric phase imaging with partially coherent light*, Phys. Rev. Lett. **80** (1998) 2586–2589.
- [23] L. de Moura, S. Ullrich, *The Lean 4 theorem prover and programming language*, in *Automated Deduction – CADE 28*, LNCS **12699** (Springer, Cham, 2021), pp. 625–635.
- [24] The mathlib Community, *The Lean mathematical library*, in *Proc. 9th ACM SIGPLAN International Conference on Certified Programs and Proofs* (ACM, New York, 2020), pp. 367–381.
- [25] J. Harrison, *Formal proof—theory and practice*, Notices Amer. Math. Soc. **55** (2008) 1395–1406.
- [26] J. Avigad, J. Harrison, *Formally verified mathematics*, Commun. ACM **57**(4) (2014) 66–75.
- [27] W. Lau, *Closed-form envelopes of the Fourier-optics kernel in aberration-corrected LEEM*, unpublished (2026).
- [28] W. Lau, *A mixed TCC–Born inverse of LEEM Fourier optics: algorithm selection*, unpublished (2026).